

QIQC Challenge Partner Handbook

GaugeForge | Partner discussion edition | 2026-10-04.1

GaugeForge is organizing the **Agentic Quantum Coding Challenge** around QIQC Bench and Quantum-Harbor. Teams build an agent harness around the same Tencent Hunyuan 4 Preview (HY4) model and demonstrate how planning, tool use and experimental decisions improve verified quantum engineering results. The challenge starts on **8 November 2026**; the final submission deadline is **10 November 2026** and Demo Day is **12 November 2026**. Registration opening time remains To be announced.

Why participate

This is an opportunity to engage developers who can turn a model into a working scientific system. Partners can help researchers discover the challenge, contribute technical judgment, or explore reproducible outputs after the event. Expected outputs include a public automated leaderboard, expert finalist review, selected success and failure case studies, and a technical report. These are planned outputs, not promises of audience size or commercial access.

The competition uses all 49 QIQC Bench tasks, with five designated public tasks. The benchmark paper also reports 17 evaluated agentic systems; that research number does not predict event participation. Quantum-Harbor provides a virtual laboratory; the event does not promise access to physical quantum hardware.

Event conditions

Item	Current position
Host	GaugeForge
Format	Global, online; Free entry (existing partner brief)
Teams	1–3 people; no academic or professional eligibility restriction
Competition model	Hunyuan 4 Preview (HY4); Tencent provider candidate; exact API ID, frozen snapshot and development allowance: To be announced
Official evaluations	Up to 3 complete benchmark evaluations per team
Final review	Automated ranking selects 5 eligible finalists; experts determine their final order
Code and data	AGPL code; organizer and participating partners have data and trajectory usage rights under disclosed terms

Partnership status

No external partner is confirmed in the current event record. Tencent is the prospective model/API support provider. QCoder is a prospective community and expert-review collaborator mentioned in an earlier brief. Neither name signifies an agreement, endorsement or logo permission. A named partnership will appear only after both sides approve the scope and public wording.

Contact **Shenhao Miao, Founder, GaugeForge** at shenhao.miao@gauge-forge.com to discuss a specific contribution. This handbook does not establish a sponsorship commitment or grant access to private evaluation materials.

Collaboration options

Contribution	Partner delivers	GaugeForge delivers	Estimated involvement
Community distribution	One approved announcement and one reminder through relevant channels	Copy, poster, public links and named contact	1–2 hours
Expert panel	One qualified reviewer, conflict disclosure, independent rubric scores	Finalist materials, rubric, briefing and review secretary	5–7 hours including preparation
Technical workshop	A focused session on agent engineering or quantum practice	Event hosting, public examples, moderation and recording consent	3–5 hours including preparation
Model/API support	Frozen HY4 access, quota terms, service contact and incident support	Integration, metering, participant documentation and fair-use enforcement	1–2 engineering days plus event on-call coverage
Compute support	Agreed worker capacity, regions and run window	Container manifests, capacity plan and operational contact	Scope agreed after load testing
Case study or report	Review of attributed commentary and permission to quote	Draft case study, factual review and final attribution	2–4 hours
Post-event experiment	A separate feasibility discussion and environment constraints	Selected reproducible harness and proposed test protocol	Separate scope; no hardware integration assumed

These are optional contributions. A community partner does not have to provide compute, judge entrants or sponsor prizes. A model partner does not automatically receive participant contact details or unrestricted private logs.

Practical agreement

Each side should name one primary contact and one backup. Before announcing a relationship, record the chosen contribution, deliverable, due date, staff time, attribution wording, data access, and withdrawal or replacement process. GaugeForge owns event operations and public rules; partners own the accuracy of their own supplied claims and their approved trademarks.

Use the supplied public poster and links without changing dates, prize claims or model wording. Logo use requires the brand owner's permission. Do not add institution names as endorsers based on an employee's participation. Co-authored case studies need factual review by both parties; publication of participant trajectories follows the participant notice and any necessary redaction.

Expert review and fairness

The complete automated score determines the five eligible finalists after reproducibility checks and appeals. Experts then determine the final order of those finalists. The automated leaderboard remains separately visible, so readers can distinguish measured performance from expert judgment.

The proposed rubric is scientific reliability 40%, harness design and supporting evidence 30%, reusability 20%, and clarity 10%. Each finalist presents a success and a failure case. At least three non-conflicted reviewers should score each team; reviewers disclose employment, mentoring, investment and relevant collaboration relationships, and recuse where needed. The organizer provides a backup reviewer. Reviewers must not request hidden tasks or answer keys, coach a team on private evaluation content, or promise an outcome.

The proposed session is 8 minutes of presentation and 7 minutes of questions per finalist, followed by independent scoring. Reviewers receive the final rubric before confirming participation. Scores and concise evidence-based reasons are returned to the review secretary; procedural appeals are resolved before awards become final.

Model and compute support

All official HY4 runs must use the same frozen API model, settings and per-task resource limits. The paper's label "Hunyuan T4*" describes a preview evaluation and is not an executable API identifier. A provider must confirm its production or event-specific model ID, version freeze, rate limits, service region, data retention and escalation route before integration.

Unlimited development credits are an organizer expectation conditional on a Tencent agreement, **not a confirmed entitlement**. The 3-run official evaluation cap remains independent of development credits. Hardware simulations, model tokens and wall time all require explicit limits. External model assistance during official execution is outside the fixed-model division.

Prizes and communication

Placement awards: 1st USD 1,500; 2nd USD 700; 3rd USD 700; 4th USD 500; 5th USD 500. The USD 15,000 Challenger Award goes only to the **first qualifying HY4 team to strictly exceed the frozen GPT-6 Astra baseline on the hidden leaderboard**. Timestamp, stacking and baseline details are to be announced. No partner is implied to fund these awards.

Checkpoint	Discussion or deliverable
Before public partnership naming	Scope, contact, attribution, logo and data permissions
Before formal event launch	Final public copy and link check; schedule confirmation
Before competition start on 8 November 2026	Technical readiness, reviewer confirmation, escalation contacts
Before Demo Day, 12 November 2026	Finalist evidence, conflicts and rubric briefing
After final awards	Approved case studies, impact summary and optional next experiment

Public references: [Benchmark](#), [Quantum-Harbor](#), [Research paper](#). Email: shenhao.miao@gauge-forge.com.

Registration link: **To be announced**; the benchmark page provides research information and is not a registration form.